

1. Introduction

Mind maps turn a wall of prose into something you can scan in one glance, which is exactly what a reader needs when the source runs to many pages. This fixture exists so the PDF extractor can be tested without shipping a multi-megabyte sample document in the repository.

2. Methods

Text is pulled page by page, then the lines are classified: numbered or short lines become headings and everything else becomes a paragraph. Bookmarks, when the document has them, override that guess entirely. The outline generator only ever sees the resulting block structure.

2.1 Limits

Scanned documents have no text layer at all, so they are rejected with a clear error rather than silently producing an empty mind map.

Optical character recognition is out of scope for this tool.